

```
# prompt: import warnings

import warnings
warnings.filterwarnings('ignore')

import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

advertising = pd.DataFrame(pd.read_csv("/content/advertising.csv"))

# prompt: advertising.head()

advertising.head()
```

	TV	Radio	Newspaper	Sales	
0	230.1	37.8	69.2	22.1	
1	44.5	39.3	45.1	10.4	
2	17.2	45.9	69.3	12.0	
3	151.5	41.3	58.5	16.5	
4	180.8	10.8	58.4	17.9	

```
# prompt: advertising.shape()

advertising.shape

(200, 4)
```

```
advertising.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    Radio       200 non-null    float64
2    Newspaper   200 non-null    float64
3    Sales       200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```


```
advertising.describe()
```

	TV	Radio	Newspaper	Sales	
count	200.000000	200.000000	200.000000	200.000000	
mean	147.042500	23.264000	30.554000	15.130500	
std	85.854236	14.846809	21.778621	5.283892	
min	0.700000	0.000000	0.300000	1.600000	
25%	74.375000	9.975000	12.750000	11.000000	
50%	149.750000	22.900000	25.750000	16.000000	
75%	218.825000	36.525000	45.100000	19.050000	
max	296.400000	49.600000	114.000000	27.000000	


```
advertising.isnull().sum()*100/advertising.shape[0]

TV          0.0
Radio       0.0
Newspaper   0.0
Sales       0.0
dtype: float64
```

```
fig, axs = plt.subplots(3, figsize = (5,5))
plt1 = sns.boxplot(advertising['TV'], ax = axs[0])
plt2 = sns.boxplot(advertising['Newspaper'], ax = axs[1])
plt3 = sns.boxplot(advertising['Radio'], ax = axs[2])
plt.tight_layout()
```



```
sns.boxplot(advertising['Sales'])
plt.show()
```



```
x = advertising['TV']

y = advertising['Sales']

from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, train_size= 0.7)

x_train.head()

24      62.3
176     248.4
40      202.5
52      216.4
94      107.4
Name: TV, dtype: float64
```

```
y_train.head()

24      9.7
176    20.2
40     16.6
52     22.6
94     11.5
Name: Sales, dtype: float64

import statsmodels.api as sm

x_train_sm = sm.add_constant(x_train)

lr = sm.OLS(y_train, x_train_sm).fit()

lr.params

const    6.960347
TV       0.054110
dtype: float64

print(lr.summary())
```

➡

OLS Regression Results						
=====						
Dep. Variable:	Sales	R-squared:	0.822			
Model:	OLS	Adj. R-squared:	0.820			
Method:	Least Squares	F-statistic:	635.3			
Date:	Wed, 31 Jan 2024	Prob (F-statistic):	1.69e-53			
Time:	09:24:28	Log-Likelihood:	-304.39			
No. Observations:	140	AIC:	612.8			
Df Residuals:	138	BIC:	618.7			
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	6.9603	0.362	19.235	0.000	6.245	7.676
TV	0.0541	0.002	25.205	0.000	0.050	0.058
=====						
Omnibus:	0.582	Durbin-Watson:	2.123			
Prob(Omnibus):	0.748	Jarque-Bera (JB):	0.245			
Skew:	0.002	Prob(JB):	0.885			
Kurtosis:	3.205	Cond. No.	337.			
=====						
Notes:						
[1] Standard Errors assume that the covariance matrix of the errors is correct						

```
y_train_pred = lr.predict(x_train_sm)
res = (y_train - y_train_pred)

x_test_sm = sm.add_constant(x_test)
y_pred = lr.predict(x_test_sm)

y_pred.head()

126    7.382404
120   14.606074
83    10.661464
134    8.957002
37    11.002356
dtype: float64

from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

np.sqrt(mean_squared_error(y_test, y_pred))

2.6498686943078926

r_squared = r2_score(y_test, y_pred)
r_squared

0.78588599097311
```